THE COMPLETE INTERVIEW PREP GUIDE

TOP 200 AI INTERVIEW QUESTIONS & ANSWERS

Artificial Intelligence | Machine Learning | Deep Learning
NLP & Generative AI | MLOps | Ethics | Latest AI Trends

10 Categories • 200 Questions • 200 Detailed Answers

COMPILED & PREPARED BY

THANGARAJ

AI/ML Intern | Data Science Developer

Coimbatore, Tamil Nadu, India

Prepared July 2026

A free-to-share resource for every AI & Data Science aspirant

TABLE OF CONTENTS

10 Categories • 200 Questions

How to use this guide: Each of the 200 questions below mirrors real interview patterns across AI, Machine Learning, Deep Learning, NLP/Generative AI, Python, MLOps, and AI Ethics. Answers are written concisely so they can be revised quickly before an interview, and are structured so any student, job-seeker, or working professional can use them — not just for AI/ML roles.

SECTION 01

ARTIFICIAL INTELLIGENCE BASICS

Questions 1-20

Q1. What is Artificial Intelligence (AI), and how does it differ from traditional programming?

AI is the field of building systems that can perform tasks requiring human-like intelligence, such as learning, reasoning, and decision-making. Traditional programming follows fixed, explicitly coded rules to produce outputs from inputs, while AI (especially Machine Learning) learns patterns from data to make predictions, allowing it to handle situations the programmer never explicitly coded for.

Q2. What are the different types of Artificial Intelligence (Narrow AI, General AI, and Super AI)?

Narrow AI (Weak AI) performs a specific task well, like spam filters or voice assistants. General AI (Strong AI) would match human-level intelligence across any task, which does not exist yet. Super AI is a hypothetical future stage where machine intelligence surpasses human intelligence in every domain.

Q3. What is the difference between Narrow AI, General AI, and Super AI?

Narrow AI is task-specific and is what exists today (e.g., ChatGPT, recommendation engines). General AI would reason and learn across any domain like a human. Super AI would exceed human cognitive ability entirely. The difference lies in scope and level of intelligence.

Q4. What is the difference between Artificial Intelligence, Machine Learning, and Deep Learning?

AI is the broad umbrella goal of making machines intelligent. Machine Learning is a subset of AI where systems learn patterns from data instead of being explicitly programmed. Deep Learning is a subset of ML that uses multi-layered neural networks to learn complex patterns, especially from unstructured data like images and text.

Q5. What are AI agents, and how do they work?

AI agents are systems that perceive their environment, reason about goals, and take actions autonomously to achieve an objective, often in a loop of observe-plan-act. Modern AI agents built on LLMs can call tools, use memory, and break tasks into sub-steps to complete multi-stage workflows with minimal human input.

Q6. What is intelligent behavior in Artificial Intelligence?

Intelligent behavior refers to a system's ability to perceive its environment, learn from experience, reason about uncertain information, and act to achieve goals effectively -- similar to how humans adapt behavior based on context and feedback.

Q7. What are expert systems, and where are they used?

Expert systems are rule-based AI programs that emulate the decision-making of a human expert using a knowledge base and an inference engine. They are used in domains like medical diagnosis, financial fraud detection, and technical troubleshooting where decisions follow well-defined expert rules.

Q8. What is knowledge representation in AI, and why is it important?

Knowledge representation is the method of encoding real-world information (facts, rules, relationships) into a format a machine can reason with, such as semantic networks, ontologies, or logic rules. It is important because how knowledge is structured directly affects an AI system's ability to reason and answer queries correctly.

Q9. What is search in Artificial Intelligence, and what are its different types?

Search in AI is the process of exploring possible states or actions to find a solution or goal. Types include uninformed search (BFS, DFS) which explores blindly, and informed/heuristic search (A*, Greedy Best-First) which uses domain knowledge to search more efficiently.

Q10. What is heuristic search, and how does it improve search efficiency?

Heuristic search uses a problem-specific estimate (heuristic function) of how close a state is to the goal to guide exploration, rather than exploring all possibilities blindly. This drastically reduces the search space and speeds up finding a solution, as seen in algorithms like A* search.

Q11. What is a state space, and how is it used in AI problem-solving?

A state space is the set of all possible configurations (states) a system can be in, connected by valid actions/transitions. AI problem-solving is framed as searching through this state space from an initial state to a goal state, e.g., in puzzles, pathfinding, or planning problems.

Q12. What is reinforcement in AI, and how does Reinforcement Learning work?

Reinforcement is the feedback (reward or penalty) an agent receives for its actions. In Reinforcement Learning, an agent interacts with an environment, takes actions, and learns a policy that maximizes cumulative reward over time through trial and error, without needing labeled data.

Q13. What are the real-world applications of Artificial Intelligence?

AI powers applications like recommendation systems (Netflix, Amazon), voice assistants (Alexa, Siri), fraud detection in banking, medical image diagnosis, self-driving cars, chatbots, predictive maintenance in manufacturing, and generative tools for content creation.

Q14. What are the major limitations and challenges of Artificial Intelligence?

Key challenges include data dependency and bias, lack of true reasoning/common sense, high computational cost, poor explainability (black-box models), vulnerability to adversarial attacks, ethical and privacy concerns, and difficulty generalizing beyond the data distribution it was trained on.

Q15. What is the Turing Test, and why is it significant?

Proposed by Alan Turing in 1950, the Turing Test evaluates whether a machine's conversational responses are indistinguishable from a human's to a human judge. It is significant as one of the earliest philosophical benchmarks for machine intelligence, even though it is now considered an imperfect measure.

Q16. What is computer vision, and what are its common applications?

Computer vision is the field of AI that enables machines to interpret and understand visual information from images or video. Common applications include facial recognition, object detection, medical imaging diagnosis, autonomous vehicle perception, and quality inspection in manufacturing.

Q17. What is Natural Language Processing (NLP), and how is it important?

NLP is the branch of AI focused on enabling machines to understand, interpret, and generate human language. It is important because it powers chatbots, translation tools, sentiment analysis, search engines, and virtual assistants that interact with people in natural language.

Q18. What is speech recognition, and how does it work?

Speech recognition converts spoken audio into text by extracting acoustic features from sound waves and mapping them to phonemes and words using models like deep neural networks (e.g., transformer-based or RNN-based acoustic models) trained on large speech datasets.

Q19. What are recommendation systems, and how do they function?

Recommendation systems predict items a user is likely to prefer, using approaches like collaborative filtering (based on similar users/items), content-based filtering (based on item attributes), or hybrid models, commonly used by platforms like Netflix, Amazon, and YouTube.

Q20. What are AI ethics, and why are they important?

AI ethics is the study of ensuring AI systems are fair, transparent, accountable, and safe. It matters because poorly governed AI can amplify bias, invade privacy, displace jobs, or make harmful decisions at scale, so ethical guardrails are essential for responsible deployment.

SECTION 02

MACHINE LEARNING FUNDAMENTALS

Questions 21-50

Q21. What is Machine Learning, and how does it work?

Machine Learning is a subset of AI where algorithms learn patterns from historical data to make predictions or decisions without being explicitly programmed for every rule. It works by fitting a model's parameters to minimize error on training data, then generalizing that learned pattern to new, unseen data.

Q22. What are the different types of Machine Learning?

The main types are Supervised Learning (learning from labeled data), Unsupervised Learning (finding patterns in unlabeled data), Reinforcement Learning (learning via rewards from an environment), and Semi-Supervised Learning (a mix of labeled and unlabeled data).

Q23. What is Supervised Learning, and when is it used?

Supervised Learning trains a model on labeled data (input-output pairs) so it learns to map inputs to correct outputs. It is used when historical labeled data is available, for tasks like classification (spam detection) and regression (price prediction).

Q24. What is Unsupervised Learning, and what are its applications?

Unsupervised Learning finds hidden patterns or structure in unlabeled data without predefined outputs. Applications include customer segmentation (clustering), anomaly detection, dimensionality reduction, and market basket analysis.

Q25. What is Reinforcement Learning, and how does it differ from other learning methods?

Reinforcement Learning trains an agent to make sequential decisions by rewarding good actions and penalizing bad ones through interaction with an environment. Unlike supervised/unsupervised learning, it does not rely on a fixed dataset -- it learns from ongoing trial-and-error feedback.

Q26. What is a dataset, and what are its components?

A dataset is a structured collection of data used to train and evaluate ML models. Its components typically include features (independent variables/input columns), labels/targets (the output to predict, in supervised learning), and rows representing individual samples/observations.

Q27. What is the difference between training data, validation data, and testing data?

Training data is used to fit the model's parameters. Validation data is used during development to tune hyperparameters and select the best model without touching the test set. Testing data is held out entirely and used only at the end to evaluate the model's real-world generalization performance.

Q28. What is the difference between features and labels in Machine Learning?

Features are the independent input variables (e.g., age, income) used to make a prediction. Labels are the dependent output values (e.g., 'churned' or 'not churned') that the model is trained to predict, present only in supervised learning.

Q29. What is overfitting, and how can it be prevented?

Overfitting happens when a model learns the training data too well, including its noise, resulting in poor performance on new data. It can be prevented using techniques like cross-validation, regularization (L1/L2), pruning, dropout, early stopping, and gathering more training data.

Q30. What is underfitting, and how can it be fixed?

Underfitting occurs when a model is too simple to capture the underlying pattern in the data, leading to poor performance on both training and test sets. It can be fixed by using a more complex model, adding relevant features, reducing regularization, or training longer.

Q31. What is the bias-variance tradeoff in Machine Learning?

Bias is error from overly simplistic assumptions (underfitting), while variance is error from excessive sensitivity to training data fluctuations (overfitting). The bias-variance tradeoff describes how reducing one tends to increase the other, and the goal is to find a model complexity that balances both for the best generalization.

Q32. What is cross-validation, and why is it important?

Cross-validation is a technique (e.g., k-fold) that splits the data into multiple train/validation folds to evaluate model performance more reliably than a single train-test split. It is important because it reduces the risk of a lucky/unlucky split skewing performance estimates and helps detect overfitting.

Q33. What is feature engineering, and why does it matter?

Feature engineering is the process of creating, transforming, or selecting input variables to improve a model's predictive performance, such as creating ratios, encoding categories, or extracting date parts. It matters because well-engineered features often improve accuracy more than switching algorithms.

Q34. What is feature scaling, and when should it be used?

Feature scaling transforms numeric features to a common scale (e.g., via normalization or standardization) so that no feature dominates due to its magnitude. It should be used for distance-based algorithms (KNN, SVM, K-Means) and gradient-based models (neural networks, logistic regression).

Q35. What is the difference between normalization and standardization?

Normalization (Min-Max scaling) rescales values to a fixed range, typically 0 to 1, preserving the shape of the original distribution. Standardization (Z-score scaling) centers data around a mean of 0 with a standard deviation of 1, and is less sensitive to outliers than normalization.

Q36. How do you evaluate the performance of a Machine Learning model?

For classification, metrics include accuracy, precision, recall, F1-score, ROC-AUC, and confusion matrix. For regression, metrics include MAE, MSE, RMSE, and R-squared. The right metric depends on the business problem, especially with imbalanced classes.

Q37. What is the difference between accuracy and precision?

Accuracy is the overall proportion of correct predictions out of all predictions. Precision is the proportion of predicted positives that are actually positive (TP / (TP+FP)). Accuracy can be misleading on imbalanced data, while precision focuses on the quality of positive predictions.

Q38. What is the difference between recall and F1-score?

Recall measures how many actual positives were correctly identified (TP / (TP+FN)), important when missing positives is costly (e.g., disease detection). F1-score is the harmonic mean of precision and recall, balancing both when there is a tradeoff between them.

Q39. What is a confusion matrix, and how do you interpret it?

A confusion matrix is a table showing True Positives, True Negatives, False Positives, and False Negatives for a classification model. It is interpreted by examining these counts to compute metrics like precision, recall, and to understand what types of errors the model is making.

Q40. What is the ROC-AUC curve, and why is it useful?

The ROC curve plots the True Positive Rate against the False Positive Rate at various classification thresholds, and AUC (Area Under the Curve) summarizes this into a single score from 0 to 1. It is useful because it evaluates a classifier's performance across all thresholds, independent of a fixed cutoff.

Q41. What is the difference between regression and classification?

Regression predicts continuous numeric outputs (e.g., house price), while classification predicts discrete categorical outputs (e.g., spam or not spam). The choice of algorithm, loss function, and evaluation metric differs based on which type of output is needed.

Q42. What is clustering, and what are the common clustering algorithms?

Clustering is an unsupervised technique that groups similar data points together based on feature similarity, without predefined labels. Common algorithms include K-Means, Hierarchical Clustering, and DBSCAN, each differing in how they define and form clusters.

Q43. What is dimensionality reduction, and why is it important?

Dimensionality reduction reduces the number of input features while preserving as much meaningful information as possible, using techniques like PCA or t-SNE. It is important for reducing computation cost, mitigating the curse of dimensionality, removing noise/redundancy, and enabling visualization of high-dimensional data.

Q44. What is Principal Component Analysis (PCA), and how does it work?

PCA is a dimensionality reduction technique that transforms correlated features into a smaller set of uncorrelated variables called principal components, ordered by the variance they explain. It works by computing the covariance matrix's eigenvectors and projecting data onto the top eigenvectors capturing the most variance.

Q45. What is hyperparameter tuning, and why is it necessary?

Hyperparameter tuning is the process of finding the best configuration values (like learning rate, tree depth, or number of clusters) that are set before training, since they are not learned from data. It is necessary because the right hyperparameters can significantly improve model performance and prevent over/underfitting.

Q46. What is Grid Search, and how does it optimize model performance?

Grid Search exhaustively tries every combination of specified hyperparameter values and evaluates each using cross-validation to find the combination that yields the best performance. It optimizes performance by systematically searching the hyperparameter space rather than guessing manually.

Q47. What is Random Search, and how does it compare to Grid Search?

Random Search samples random combinations of hyperparameters from specified ranges rather than trying every combination. Compared to Grid Search, it is often more efficient for large search spaces because it can find good configurations faster with fewer evaluations, though it isn't exhaustive.

Q48. What is AutoML, and what are its benefits?

AutoML automates the end-to-end process of model selection, feature engineering, and hyperparameter tuning with minimal human intervention. Its benefits include faster experimentation, democratizing ML for non-experts, and often discovering strong models that would take much longer to find manually.

Q49. What is ensemble learning, and why is it effective?

Ensemble learning combines predictions from multiple models to produce a stronger overall prediction than any single model. It is effective because it reduces variance (bagging), reduces bias (boosting), and averages out individual model errors, improving robustness and accuracy.

Q50. What is the difference between bagging and boosting?

Bagging (e.g., Random Forest) trains multiple models independently and in parallel on random data subsets, then averages/votes their outputs to reduce variance. Boosting (e.g., XGBoost) trains models sequentially, where each new model focuses on correcting the errors of the previous one, reducing bias.

SECTION 03

DEEP LEARNING

Questions 51-80

Q51. What is Deep Learning, and how is it different from Machine Learning?

Deep Learning is a subset of Machine Learning that uses multi-layered artificial neural networks to automatically learn hierarchical feature representations from raw data. Unlike traditional ML, which often needs manual feature engineering, Deep Learning learns features directly, especially excelling with unstructured data like images, audio, and text.

Q52. What is an artificial neural network, and how does it work?

An artificial neural network is a computational model inspired by the brain, made of layers of interconnected neurons (nodes) with weighted connections. It works by passing input data through layers where each neuron applies a weighted sum and activation function, adjusting weights via backpropagation to minimize prediction error.

Q53. What is a perceptron in Deep Learning?

A perceptron is the simplest type of artificial neuron -- it takes weighted inputs, sums them, adds a bias, and passes the result through an activation function to produce a binary output. It forms the basic building block of larger neural networks.

Q54. What are hidden layers in a neural network, and why are they used?

Hidden layers are the layers between the input and output layers where the network learns intermediate feature representations. They are used because they allow the network to model complex, non-linear relationships in data that a single layer could not capture.

Q55. What are activation functions, and why are they used?

Activation functions introduce non-linearity into a neural network, enabling it to learn complex patterns rather than just linear relationships. Without them, a neural network would behave like a simple linear model regardless of depth.

Q56. What is the difference between ReLU, Sigmoid, and Tanh activation functions?

ReLU outputs the input directly if positive and zero otherwise, is computationally efficient, and is widely used in hidden layers. Sigmoid squashes outputs between 0 and 1, useful for binary probabilities but prone to vanishing gradients. Tanh squashes outputs between -1 and 1, centered at zero, but also suffers from vanishing gradients for large inputs.

Q57. What is backpropagation, and how does it train neural networks?

Backpropagation is the algorithm used to train neural networks by computing the gradient of the loss function with respect to each weight, using the chain rule, and propagating this error backward from the output layer to the input layer. These gradients are then used by an optimizer to update the weights and reduce error.

Q58. What is gradient descent, and what are its different variants?

Gradient descent is an optimization algorithm that iteratively updates model parameters in the direction that reduces the loss function, using the gradient. Variants include Batch Gradient Descent (uses full dataset per step), Stochastic Gradient Descent (uses one sample per step), and Mini-Batch Gradient Descent (uses small batches, the most common in practice).

Q59. What is an optimizer, and what are the common optimization algorithms?

An optimizer is the algorithm that updates a neural network's weights to minimize the loss function based on computed gradients. Common optimizers include SGD, Momentum, RMSprop, and Adam, with Adam being the most widely used due to its adaptive learning rates and fast convergence.

Q60. What is a loss function, and why is it important?

A loss function quantifies the difference between a model's predictions and the actual target values. It is important because it is the objective the training process minimizes -- the choice of loss function (e.g., cross-entropy vs. MSE) directly shapes what the model learns to optimize for.

Q61. What is batch size, and how does it affect training?

Batch size is the number of training samples processed before the model's weights are updated. A smaller batch size gives noisier but more frequent updates and can generalize better, while a larger batch size gives smoother, faster (per epoch) updates but requires more memory and can converge to sharper minima.

Q62. What are epochs, and how many are typically required?

An epoch is one complete pass of the entire training dataset through the neural network. The number required varies widely (from a handful to hundreds) depending on dataset size, model complexity, and convergence behavior, and is typically tuned using validation loss and early stopping rather than a fixed number.

Q63. What is dropout, and how does it prevent overfitting?

Dropout is a regularization technique that randomly deactivates a fraction of neurons during each training iteration. It prevents overfitting by forcing the network to not rely too heavily on any single neuron, encouraging more robust, distributed feature representations.

Q64. What is batch normalization, and why is it used?

Batch normalization normalizes the inputs to each layer within a mini-batch to have a stable mean and variance during training. It is used because it speeds up training, allows higher learning rates, reduces sensitivity to weight initialization, and has a mild regularizing effect.

Q65. What are Convolutional Neural Networks (CNNs), and where are they used?

CNNs are deep learning architectures that use convolutional filters to automatically extract spatial features like edges, textures, and shapes from grid-like data. They are primarily used for image and video tasks such as image classification, object detection, and facial recognition.

Q66. What are the common applications of CNNs?

CNNs are commonly used for image classification, object detection, facial recognition, medical image analysis (e.g., tumor detection), self-driving car perception, and even some NLP and time-series tasks using 1D convolutions.

Q67. What are Recurrent Neural Networks (RNNs), and how do they work?

RNNs are neural networks designed for sequential data, where the output from a previous time step is fed back as input to the current step, giving the network a form of memory. This allows RNNs to model dependencies over time, though they struggle with long sequences due to vanishing gradients.

Q68. What are Long Short-Term Memory (LSTM) networks?

LSTMs are a type of RNN designed to solve the vanishing gradient problem using gated memory cells (input, forget, and output gates) that control what information to retain or discard over long sequences, making them effective for long-range dependency tasks like language modeling.

Q69. What are Gated Recurrent Units (GRUs), and how do they differ from LSTMs?

GRUs are a simplified variant of LSTM that combine the forget and input gates into a single 'update gate' and merge the cell and hidden states. They differ from LSTMs by having fewer parameters and being computationally faster, while often achieving comparable performance.

Q70. What are Transformer models, and why are they important?

Transformers are deep learning architectures that use self-attention mechanisms to process entire sequences in parallel, rather than sequentially like RNNs. They are important because they enable much faster training on long sequences and form the foundation of modern LLMs like GPT and BERT.

Q71. What is self-attention in Transformer models?

Self-attention is a mechanism that allows each token in a sequence to weigh the relevance of every other token when computing its representation, capturing contextual relationships regardless of their distance in the sequence.

Q72. What is the attention mechanism, and how does it work?

The attention mechanism computes a weighted combination of input elements, where the weights (attention scores) reflect how relevant each element is to the current context, typically calculated via Query, Key, and Value vectors and a softmax-normalized similarity score.

Q73. What is transfer learning, and when should it be used?

Transfer learning reuses a model pretrained on a large dataset (e.g., ImageNet or a large text corpus) and fine-tunes it for a new, related task with a smaller dataset. It should be used when labeled data is limited or when training from scratch would be too costly or time-consuming.

Q74. What are embeddings in Deep Learning?

Embeddings are dense, low-dimensional vector representations of data (words, images, users, etc.) that capture semantic meaning and relationships, such that similar items are placed closer together in the vector space.

Q75. What is fine-tuning, and why is it useful?

Fine-tuning is the process of taking a pretrained model and continuing training on a smaller, task-specific dataset to adapt it to a new task. It is useful because it leverages the general knowledge already learned, achieving strong performance with far less data and compute than training from scratch.

Q76. What is multimodal AI, and how does it work?

Multimodal AI processes and relates multiple types of data (text, images, audio, video) within a single model. It works by encoding each modality into a shared representation space so the model can reason jointly across modalities, e.g., answering questions about an image.

Q77. What are Generative Adversarial Networks (GANs)?

GANs consist of two neural networks -- a generator that creates fake data and a discriminator that tries to distinguish fake from real data -- trained in competition. This adversarial process pushes the generator to produce increasingly realistic synthetic data, such as images.

Q78. What are Autoencoders, and what are their applications?

Autoencoders are neural networks trained to compress input data into a smaller latent representation (encoder) and then reconstruct the original input from it (decoder). Applications include dimensionality reduction, anomaly detection, denoising, and generative modeling (variational autoencoders).

Q79. What are diffusion models, and how do they generate images?

Diffusion models generate images by learning to reverse a gradual noising process -- starting from pure random noise, the model iteratively denoises it step by step, guided by learned patterns, until a coherent image emerges. They power tools like Stable Diffusion and DALL-E.

Q80. What are foundation models, and why are they significant?

Foundation models are large models pretrained on massive, broad datasets (e.g., GPT, BERT) that can be adapted to many downstream tasks via fine-tuning or prompting. They are significant because they shift AI development from building task-specific models to adapting general-purpose ones, greatly reducing development time and cost.

SECTION 04

NLP & GENERATIVE AI

Questions 81-110

Q81. What is Natural Language Processing (NLP), and what are its applications?

NLP is the field of AI focused on enabling computers to understand, interpret, and generate human language. Applications include chatbots, machine translation, sentiment analysis, text summarization, search engines, and voice assistants.

Q82. What are the main stages of the NLP pipeline?

A typical NLP pipeline includes text preprocessing (cleaning, tokenization, stopword removal), feature extraction/representation (TF-IDF, embeddings), model training or inference, and post-processing/evaluation of the output.

Q83. What is tokenization, and why is it important?

Tokenization is the process of splitting text into smaller units like words, subwords, or characters (tokens). It is important because it converts raw text into discrete units that models can numerically process and analyze.

Q84. What is the difference between stemming and lemmatization?

Stemming crudely chops word endings using fixed rules to get a root form (e.g., 'studies' to 'studi'), which may not be a real word. Lemmatization uses vocabulary and grammar rules to reduce a word to its proper dictionary base form (e.g., 'studies' to 'study'), producing more linguistically accurate results.

Q85. What are stop words, and why are they removed?

Stop words are common, low-information words like 'the', 'is', and 'and'. They are often removed in traditional NLP pipelines to reduce noise and dimensionality, focusing analysis on more meaningful, content-bearing words -- though modern deep learning models often keep them since context matters.

Q86. What is TF-IDF, and how is it calculated?

TF-IDF (Term Frequency-Inverse Document Frequency) measures how important a word is to a document within a corpus. It is calculated by multiplying Term Frequency (how often a word appears in a document) by Inverse Document Frequency (a measure that down-weights words common across many documents).

Q87. What is Word2Vec, and how does it generate word embeddings?

Word2Vec is a neural network-based technique that learns dense word embeddings by predicting a word from its context (CBOW) or predicting context words from a target word (Skip-gram), resulting in vectors where semantically similar words are close together.

Q88. What is GloVe, and how does it differ from Word2Vec?

GloVe (Global Vectors) generates word embeddings by factorizing a global word co-occurrence matrix built from the entire corpus, capturing overall statistical patterns. Unlike Word2Vec, which learns from local context windows, GloVe explicitly leverages global corpus statistics.

Q89. What is BERT, and how does it work?

BERT (Bidirectional Encoder Representations from Transformers) is a transformer-based language model that reads text bidirectionally (both left and right context) to build deep contextual understanding, trained using masked language modeling and next sentence prediction.

Q90. What is GPT, and how is it different from BERT?

GPT (Generative Pretrained Transformer) is a transformer decoder-only model trained to predict the next word in a sequence, making it strong at text generation. BERT is encoder-only and bidirectional, making it stronger at understanding tasks like classification, rather than generating free-form text.

Q91. What is prompt engineering, and why is it important?

Prompt engineering is the practice of crafting inputs (prompts) to guide a language model toward producing the desired output. It is important because the quality, structure, and clarity of a prompt can dramatically affect the accuracy and usefulness of an LLM's response, without needing to retrain the model.

Q92. What is zero-shot prompting, and when is it used?

Zero-shot prompting asks a model to perform a task with no prior examples given in the prompt, relying entirely on the model's pretrained knowledge. It is used for simple, well-understood tasks where the model's general training is sufficient to respond correctly.

Q93. What is one-shot prompting, and how does it work?

One-shot prompting provides exactly one example of the task within the prompt before asking the model to complete a new instance. It works by giving the model a pattern or format to follow, improving accuracy over zero-shot for moderately complex or ambiguous tasks.

Q94. What is few-shot prompting, and why is it effective?

Few-shot prompting provides several examples of the task in the prompt to demonstrate the expected pattern before the actual query. It is effective because it helps the model infer the desired format, tone, and reasoning style, improving performance on more complex or nuanced tasks.

Q95. What is Chain of Thought (CoT) prompting?

Chain of Thought prompting encourages a language model to reason step-by-step before arriving at a final answer, often by explicitly asking it to 'think step by step.' This improves performance on tasks requiring multi-step reasoning, like math or logic problems.

Q96. What are hallucinations in Large Language Models (LLMs)?

Hallucinations occur when an LLM generates confident, plausible-sounding text that is factually incorrect or not grounded in any real source. They happen because LLMs generate text based on learned statistical patterns, not verified facts, and are a key challenge for reliable AI deployment.

Q97. What is Retrieval-Augmented Generation (RAG), and how does it work?

RAG combines an LLM with an external knowledge retrieval step: relevant documents are retrieved from a knowledge base (often via vector similarity search) and injected into the prompt as context before the model generates its answer. This reduces hallucinations and allows the model to use up-to-date or domain-specific information beyond its training data.

Q98. What is a vector database, and why is it used in AI?

A vector database stores and indexes high-dimensional embedding vectors and enables fast similarity search (e.g., nearest neighbor search) between them. It is used in AI, particularly RAG systems, to quickly retrieve semantically relevant documents or items based on a query embedding.

Q99. What is semantic search, and how does it differ from keyword search?

Semantic search retrieves results based on the meaning and context of a query, using embeddings to find conceptually similar content even if exact words differ. Keyword search, by contrast, matches exact or partial text strings, missing results that are relevant in meaning but phrased differently.

Q100. What are embeddings, and why are they important in NLP?

Embeddings are dense numeric vector representations of words, sentences, or documents that capture semantic meaning, such that similar meanings map to nearby vectors. They are important because they let models perform mathematical operations (like similarity search) on language, which is otherwise unstructured.

Q101. What is the context window in Large Language Models?

The context window is the maximum number of tokens (words/subwords) an LLM can process at once, including both the prompt and its generated response. Content beyond this limit is truncated or ignored, which affects how much information a model can 'remember' within a single interaction.

Q102. What are Large Language Models (LLMs), and how do they work?

LLMs are deep neural networks, typically transformer-based, trained on massive text corpora to predict and generate coherent, contextually relevant language. They work by learning statistical patterns of language during pretraining, then generating text token-by-token based on the preceding context.

Q103. What is the difference between fine-tuning and prompt engineering?

Fine-tuning updates a model's internal weights using additional training data to permanently adapt it to a task. Prompt engineering, by contrast, changes only the input given to a frozen, pretrained model to guide its output, requiring no retraining or weight updates.

Q104. What are the temperature parameter in LLMs?

Temperature controls the randomness of an LLM's output by scaling the probability distribution over the next token before sampling. Low temperature (near 0) makes output more deterministic and focused, while high temperature increases diversity and creativity, but also the risk of incoherence.

Q105. What is top-p (nucleus) sampling in text generation?

Top-p sampling selects the next token from the smallest set of tokens whose cumulative probability exceeds a threshold p, rather than always picking the single most likely token. This balances diversity and coherence better than fixed top-k sampling in many generation scenarios.

Q106. What are AI agents, and how do they interact with LLMs?

AI agents built on LLMs use the language model as a reasoning engine that plans actions, calls external tools or APIs, observes results, and iterates toward completing a goal, rather than just generating a single text response.

Q107. What is function calling in LLMs?

Function calling allows an LLM to output a structured request to invoke a specific external function or API (with defined parameters) based on the user's query, enabling the model to take real actions or fetch live data beyond its own knowledge.

Q108. What is tool calling in LLM-based applications?

Tool calling is the broader mechanism where an LLM decides which external tool (search, calculator, database, code executor, etc.) to use for a given task, formats the correct input for it, and incorporates the tool's output back into its response.

Q109. What is Model Context Protocol (MCP), and why is it useful?

MCP is a standardized protocol that allows LLM applications to connect to external data sources and tools in a consistent way, rather than needing custom integration code for every service. It is useful because it simplifies building AI agents that can securely access files, APIs, and services across different platforms.

Q110. How is AI safety implemented in Large Language Models?

AI safety in LLMs is implemented through techniques like reinforcement learning from human feedback (RLHF) for alignment, content filtering and moderation layers, red-teaming to find vulnerabilities, guardrails against harmful outputs, and constitutional or rule-based training approaches to shape model behavior.

SECTION 05

PYTHON FOR AI

Questions 111-130

Q111. Why is Python the most popular programming language for AI?

Python is popular for AI because of its simple, readable syntax, a vast ecosystem of ML/AI libraries (NumPy, Pandas, TensorFlow, PyTorch), strong community support, and easy integration with other tools, which speeds up experimentation and development.

Q112. What is NumPy, and why is it important for AI?

NumPy is a Python library for fast numerical computation using multi-dimensional arrays. It is important for AI because most ML and deep learning libraries build on NumPy's array operations, which are far faster than native Python lists for mathematical computations.

Q113. What is Pandas, and how is it used in data analysis?

Pandas is a Python library providing DataFrame and Series structures for handling tabular data. It is used for data cleaning, transformation, filtering, aggregation, merging datasets, and exploratory data analysis before feeding data into ML models.

Q114. What is a DataFrame in Pandas?

A DataFrame is a two-dimensional, labeled data structure in Pandas, similar to a spreadsheet or SQL table, with rows and columns that can hold different data types, and supports operations like filtering, grouping, and joining.

Q115. What is Matplotlib, and how is it used for visualization?

Matplotlib is a Python plotting library used to create static, animated, and interactive visualizations like line charts, bar charts, histograms, and scatter plots, commonly used to explore data distributions and visualize model results.

Q116. What is Seaborn, and how does it differ from Matplotlib?

Seaborn is a statistical visualization library built on top of Matplotlib that provides higher-level, aesthetically polished plotting functions (like heatmaps, pair plots, and violin plots) with less code, making it easier to visualize relationships in data compared to raw Matplotlib.

Q117. What is Scikit-learn, and what are its main features?

Scikit-learn is a Python library providing simple, consistent APIs for classical machine learning algorithms, including classification, regression, clustering, dimensionality reduction, model selection tools (train_test_split, GridSearchCV), and preprocessing utilities.

Q118. What is TensorFlow, and when should you use it?

TensorFlow is an open-source deep learning framework developed by Google for building and deploying neural networks at scale. It should be used when you need production-grade deployment support, mobile/edge deployment (TensorFlow Lite), or large-scale distributed training.

Q119. What is PyTorch, and why is it popular?

PyTorch is an open-source deep learning framework developed by Meta, known for its dynamic computation graph, intuitive Pythonic syntax, and ease of debugging. It is popular in research and increasingly in production because it feels natural to Python developers and offers great flexibility.

Q120. What is Keras, and how does it simplify Deep Learning?

Keras is a high-level neural network API that runs on top of TensorFlow, simplifying deep learning by providing simple, readable functions to build, train, and evaluate models with far less boilerplate code than working directly with low-level tensor operations.

Q121. What is OpenCV, and what are its applications?

OpenCV is an open-source computer vision library offering tools for image and video processing. Applications include object detection, face recognition, image filtering, edge detection, motion tracking, and augmented reality.

Q122. What is the Hugging Face Transformers library?

Hugging Face Transformers is a Python library providing easy access to thousands of pretrained transformer models (like BERT, GPT, T5) for tasks such as text classification, generation, translation, and question answering, with simple APIs for fine-tuning and inference.

Q123. What is LangChain, and how is it used in LLM applications?

LangChain is a framework for building applications powered by LLMs by chaining together prompts, tools, memory, and data retrieval steps. It is used to build chatbots, RAG pipelines, and AI agents that combine LLM reasoning with external data and tools.

Q124. What is LlamaIndex, and what problem does it solve?

LlamaIndex is a framework focused on connecting LLMs to external/private data sources by building indexes (often vector-based) over documents. It solves the problem of efficiently retrieving relevant context from large document collections to feed into an LLM for accurate, grounded responses.

Q125. What is Ollama, and how is it used for running local LLMs?

Ollama is a tool that lets you download, run, and manage open-source LLMs locally on your own machine with a simple command-line interface. It is used to run models like Llama or Mistral offline, useful for privacy, cost savings, and experimentation without API calls.

Q126. How do you use the OpenAI API in AI applications?

The OpenAI API is used by sending HTTP requests with a prompt/messages payload and an API key to OpenAI's endpoints (e.g., chat completions), which return generated text, embeddings, or other outputs that can be integrated into applications like chatbots or content generators.

Q127. How do you use the Anthropic API for LLM development?

The Anthropic API works similarly, sending a request with model name, messages, and parameters to Anthropic's /v1/messages endpoint, which returns Claude's response, supporting features like tool use, system prompts, and streaming for building AI-powered applications.

Q128. How do you use the Google Gemini API in AI projects?

The Google Gemini API is accessed via Google's AI Studio or Vertex AI, where you send prompts (text, image, or multimodal) to Gemini models through an API key or Google Cloud credentials, receiving generated responses for use in applications.

Q129. What is MLflow, and why is it important in MLOps?

MLflow is an open-source platform for managing the ML lifecycle, including experiment tracking, model packaging, a model registry, and deployment. It is important in MLOps because it brings reproducibility and organization to what can otherwise be a chaotic experimentation process.

Q130. What is Weights & Biases, and how is it used for experiment tracking?

Weights & Biases (W&B;) is a platform for tracking ML experiments, logging metrics, visualizing training curves, comparing model runs, and managing datasets/models. It is used to keep experiments organized and reproducible, especially when running many training iterations with different hyperparameters.

SECTION 06

DATA SCIENCE & MODEL EVALUATION

Questions 131-150

Q131. What is data preprocessing, and why is it important?

Data preprocessing is the step of cleaning and transforming raw data (handling missing values, encoding categories, scaling features) into a format suitable for modeling. It is important because real-world data is messy, and model quality is highly dependent on the quality of the input data ('garbage in, garbage out').

Q132. How do you handle missing values in a dataset?

Missing values can be handled by removing rows/columns with excessive missingness, imputing with statistics (mean, median, mode), using model-based imputation (KNN imputer), or using algorithms that natively handle missing data, depending on how much data is missing and why.

Q133. What are outliers, and how do you detect and handle them?

Outliers are data points that deviate significantly from the rest of the data. They can be detected using methods like the IQR rule, Z-scores, or visualization (box plots), and handled by removing them, capping/winsorizing, transforming the data, or using robust models less sensitive to outliers.

Q134. How do you encode categorical variables for Machine Learning?

Categorical variables can be encoded using One-Hot Encoding (creates binary columns per category, good for nominal data), Label Encoding (assigns integers, good for ordinal data), or advanced methods like target encoding for high-cardinality features.

Q135. What is a train-test split, and why is it required?

A train-test split divides a dataset into a portion used to train the model and a separate portion held out to evaluate it. It is required to measure how well a model generalizes to unseen data, avoiding overly optimistic performance estimates from evaluating on training data.

Q136. What is K-Fold Cross-Validation, and how does it work?

K-Fold Cross-Validation splits the data into k equal folds; the model is trained on k-1 folds and validated on the remaining fold, repeated k times so each fold serves as the validation set once. The final performance is the average across all k runs, giving a more reliable estimate than a single split.

Q137. What is feature selection, and why is it important?

Feature selection is the process of choosing the most relevant features and discarding irrelevant or redundant ones, using methods like correlation analysis, recursive feature elimination, or feature importance from tree models. It is important because it reduces overfitting, speeds up training, and improves model interpretability.

Q138. How is correlation measured, and how is it interpreted?

Correlation is commonly measured using Pearson's correlation coefficient, ranging from -1 to +1, where values near +1 indicate a strong positive relationship, near -1 indicate a strong negative relationship, and near 0 indicate little to no linear relationship between two variables.

Q139. How do you interpret a Precision-Recall Curve, and when should it be used?

A Precision-Recall curve plots precision against recall at different classification thresholds. It should be used instead of ROC-AUC when dealing with highly imbalanced datasets, as it better reflects performance on the minority (positive) class of interest.

Q140. What is Mean Squared Error (MSE), and how is it calculated?

MSE is a regression evaluation metric that measures the average of the squared differences between predicted and actual values. It is calculated as the mean of (actual - predicted) squared across all data points, penalizing larger errors more heavily due to the squaring.

Q141. What is the difference between MAE and RMSE?

MAE (Mean Absolute Error) averages the absolute differences between predictions and actual values, treating all errors equally. RMSE (Root Mean Squared Error) squares errors before averaging and then takes the square root, which penalizes larger errors more heavily than MAE.

Q142. What is the R-squared (R2) Score, and what does it indicate?

R-squared indicates the proportion of variance in the target variable that is explained by the model's features, ranging typically from 0 to 1 (higher is better). An R2 of 0.85, for example, means the model explains 85% of the variability in the outcome.

Q143. What is the Silhouette Score in clustering?

The Silhouette Score measures how well-separated and cohesive clusters are, ranging from -1 to +1, by comparing a point's average distance to points in its own cluster versus the nearest other cluster. Higher values indicate better-defined, more distinct clusters.

Q144. What is A/B Testing, and how is it used in AI?

A/B Testing is a controlled experiment where two variants (A and B) -- such as two model versions -- are shown to different user segments to statistically compare their performance on a key metric, commonly used to validate whether a new ML model actually improves real-world outcomes before full rollout.

Q145. What is explainability in Artificial Intelligence?

Explainability refers to the degree to which a model's internal decision-making process can be understood by humans. It matters for trust, debugging, regulatory compliance, and fairness auditing, especially for high-stakes decisions like credit approval or medical diagnosis.

Q146. What is SHAP, and how does it explain model predictions?

SHAP (SHapley Additive exPlanations) is a game-theory-based method that assigns each feature a contribution value showing how much it pushed a specific prediction higher or lower than the average prediction, providing both global and local model interpretability.

Q147. What is LIME, and how does it improve model interpretability?

LIME (Local Interpretable Model-agnostic Explanations) explains individual predictions by approximating the complex model locally with a simple, interpretable model (like linear regression) around that specific data point, showing which features most influenced that single prediction.

Q148. What is model drift, and how can it be detected?

Model drift is the degradation of a model's performance over time as real-world data patterns change (data drift) or the relationship between inputs and outputs changes (concept drift). It can be detected by continuously monitoring prediction accuracy, input feature distributions, and statistical tests comparing production data to training data.

Q149. What is data drift, and how does it affect AI models?

Data drift occurs when the statistical properties of input data in production shift away from the data the model was trained on. It affects AI models by silently degrading prediction accuracy, since the model's learned patterns no longer match the current data distribution.

Q150. How do you optimize the inference speed of a Deep Learning model?

Inference speed can be optimized through model quantization (reducing precision, e.g., FP32 to INT8), pruning unnecessary weights, knowledge distillation into a smaller model, using optimized runtimes (ONNX, TensorRT), batching requests, and deploying on hardware accelerators like GPUs or TPUs.

SECTION 07

AI DEPLOYMENT & MLOPS

Questions 151-170

Q151. What is MLOps, and why is it important?

MLOps is a set of practices that combines Machine Learning, DevOps, and data engineering to reliably build, deploy, monitor, and maintain ML models in production. It is important because it bridges the gap between experimental notebooks and robust, scalable, reproducible production systems.

Q152. What are the different methods for deploying Machine Learning models?

Models can be deployed via REST APIs (Flask/FastAPI), batch prediction pipelines, embedding directly into applications, serverless functions, containerized microservices (Docker/Kubernetes), or edge deployment on devices, depending on latency and scale requirements.

Q153. How do you expose an ML model using a REST API?

You expose an ML model as a REST API by wrapping the trained model's predict function inside a web framework (like Flask or FastAPI), defining an endpoint (e.g., /predict) that accepts input data as JSON, runs inference, and returns the prediction as a JSON response.

Q154. What is Docker, and why is it used in AI deployment?

Docker is a containerization platform that packages an application with all its dependencies (libraries, runtime, OS-level requirements) into a portable container. It is used in AI deployment to ensure models run consistently across different environments, avoiding 'it works on my machine' issues.

Q155. What is Kubernetes, and how does it help scale AI applications?

Kubernetes is a container orchestration platform that automates deployment, scaling, and management of containerized applications. It helps scale AI applications by automatically adding or removing model-serving instances based on traffic load, handling failover, and managing rolling updates.

Q156. What is FastAPI, and why is it popular for AI applications?

FastAPI is a modern Python web framework for building APIs, known for its high performance, automatic interactive documentation, and native support for data validation via type hints. It is popular for AI applications because it makes serving ML models as APIs fast and simple to build and test.

Q157. What is Flask, and how is it used for model deployment?

Flask is a lightweight Python web framework used to build simple REST APIs by defining routes that load a trained model and return predictions on incoming requests, commonly used for quick prototyping and smaller-scale ML deployments.

Q158. How is MLflow used in the MLOps lifecycle?

MLflow is used across the MLOps lifecycle to track experiments (parameters, metrics, artifacts), package models in a standard format, register and version models in a central model registry, and deploy models to various serving platforms.

Q159. What is model versioning, and why is it necessary?

Model versioning is the practice of tracking different iterations of a trained model along with their associated code, data, and hyperparameters. It is necessary to enable rollback to a previous stable model, reproduce results, and audit which model version made a specific prediction.

Q160. What is CI/CD, and how is it applied in AI projects?

CI/CD (Continuous Integration/Continuous Deployment) automates testing and deployment of code changes. In AI projects, it is extended to include automated model retraining, validation against performance thresholds, and safe rollout of new model versions into production pipelines.

Q161. What is a Feature Store, and what problem does it solve?

A Feature Store is a centralized repository for storing, managing, and serving ML features consistently across training and production (real-time) environments. It solves the problem of training-serving skew, where features computed differently in training versus production cause performance discrepancies.

Q162. What is a Model Registry in MLOps?

A Model Registry is a centralized system for storing, versioning, and managing the lifecycle stage (staging, production, archived) of trained ML models, enabling teams to track lineage and control which model version is currently deployed.

Q163. What is batch inference, and when should it be used?

Batch inference generates predictions for a large set of inputs all at once on a scheduled basis, rather than responding to individual real-time requests. It should be used when predictions don't need to be immediate, such as generating daily recommendation lists or nightly risk scores.

Q164. What is real-time inference, and how does it work?

Real-time inference generates predictions instantly on demand, typically through a low-latency API call, as soon as a request arrives. It works by keeping the model loaded in memory on a serving infrastructure that can respond within milliseconds to seconds.

Q165. How do you monitor AI models in production?

AI models in production are monitored by tracking prediction accuracy/performance metrics over time, input data distribution shifts (data drift), latency and error rates, and business KPIs, often using dashboards and automated alerting for anomalies.

Q166. What is drift detection, and why is it important?

Drift detection is the process of statistically monitoring whether incoming production data or model performance has shifted from what the model was trained on. It is important because undetected drift silently erodes model accuracy, leading to poor decisions without an obvious error.

Q167. What is the difference between a GPU and a CPU for AI workloads?

A CPU has fewer, more powerful cores optimized for sequential, general-purpose tasks, while a GPU has thousands of smaller cores optimized for massively parallel computation. AI workloads like matrix multiplications in neural network training benefit greatly from a GPU's parallelism, making training much faster.

Q168. What is a TPU, and when should it be used?

A TPU (Tensor Processing Unit) is a specialized hardware accelerator designed by Google specifically for tensor operations used in deep learning. It should be used for large-scale training or inference of deep learning models, particularly within the Google Cloud/TensorFlow ecosystem, where it can offer speed and cost advantages over GPUs.

Q169. What is Edge AI, and what are its advantages?

Edge AI refers to running AI models directly on local devices (phones, IoT sensors, cameras) rather than in the cloud. Its advantages include lower latency, reduced bandwidth/cost, improved privacy (data stays on-device), and the ability to function without a constant internet connection.

Q170. How do you design AI systems for scalability?

Designing AI systems for scalability involves using containerized, stateless model-serving architectures, load balancing across multiple instances, caching frequent predictions, using asynchronous/batch processing where possible, and choosing infrastructure (Kubernetes, autoscaling groups) that can grow with demand.

SECTION 08

AI ETHICS & SECURITY

Questions 171-185

Q171. What is AI bias, and how can it be reduced?

AI bias occurs when a model produces systematically unfair outcomes for certain groups, often because of imbalanced or historically biased training data. It can be reduced through diverse and representative datasets, bias auditing tools, fairness-aware algorithms, and human review of high-stakes decisions.

Q172. What is fairness in Artificial Intelligence?

Fairness in AI means ensuring a model's decisions do not systematically disadvantage individuals or groups based on sensitive attributes like race, gender, or age, often measured through statistical fairness metrics such as demographic parity or equal opportunity.

Q173. What is Responsible AI, and why is it important?

Responsible AI is the practice of designing, developing, and deploying AI systems that are fair, transparent, safe, accountable, and privacy-respecting. It is important because AI increasingly influences high-stakes decisions in hiring, lending, healthcare, and law enforcement, where harm can be significant if not carefully managed.

Q174. What is Explainable AI (XAI), and why is it needed?

Explainable AI encompasses techniques (like SHAP and LIME) that make a model's predictions understandable to humans. It is needed for building trust, meeting regulatory requirements, debugging model errors, and ensuring accountability, especially for black-box models like deep neural networks.

Q175. What is AI governance, and what are its key principles?

AI governance is the framework of policies, standards, and oversight structures that guide how AI is developed and used responsibly within an organization or society. Key principles include transparency, accountability, fairness, safety, human oversight, and compliance with legal/regulatory requirements.

Q176. Why is data privacy important in AI systems?

Data privacy is important in AI systems because models are often trained on sensitive personal data, and mishandling it can lead to identity theft, discrimination, or loss of user trust; regulations like GDPR also impose legal obligations on how such data can be collected, stored, and used.

Q177. What is differential privacy, and how does it protect user data?

Differential privacy is a mathematical technique that adds calibrated statistical noise to data or query results, ensuring that the presence or absence of any single individual's data does not significantly affect the output, thereby protecting individual privacy while still allowing useful aggregate analysis.

Q178. What are adversarial attacks in Machine Learning?

Adversarial attacks are deliberately crafted, often subtle, input perturbations designed to fool an ML model into making an incorrect prediction, such as slightly altered images that cause an image classifier to misidentify an object while looking unchanged to humans.

Q179. What is prompt injection, and how can it be prevented?

Prompt injection is an attack where malicious instructions are embedded in input text (or external content an LLM reads) to manipulate the model into ignoring its original instructions or leaking sensitive information. It can be mitigated through input sanitization, strict system prompt boundaries, output filtering, and treating untrusted content as data rather than instructions.

Q180. What is jailbreaking in Large Language Models?

Jailbreaking refers to techniques used to manipulate an LLM into bypassing its built-in safety guidelines and producing restricted or harmful content, often through clever prompt framing, roleplay scenarios, or exploiting model reasoning gaps.

Q181. What is model poisoning, and how does it affect AI systems?

Model poisoning is an attack where an adversary injects malicious or corrupted data into a model's training set to intentionally degrade its performance or embed hidden malicious behavior (a backdoor) that activates under specific conditions.

Q182. What are the copyright and intellectual property issues in AI?

Key copyright/IP issues in AI include whether training on copyrighted data without permission constitutes infringement, who owns AI-generated content, and how to prevent models from reproducing memorized copyrighted text, image, or code verbatim -- an area still evolving legally worldwide.

Q183. What are deepfakes, and what risks do they pose?

Deepfakes are AI-generated synthetic media (images, audio, or video) that convincingly depict real people saying or doing things they never did. They pose risks including misinformation, fraud, non-consensual content, reputational harm, and erosion of trust in genuine media.

Q184. What are the major AI regulations around the world?

Major AI regulatory efforts include the EU AI Act (a risk-based regulatory framework), the US Executive Orders and NIST AI Risk Management Framework, China's AI governance regulations for generative AI, and various sector-specific rules (e.g., in finance and healthcare) governing automated decision-making.

Q185. What are the ethical challenges of Generative AI?

Ethical challenges of Generative AI include misinformation and deepfakes, copyright infringement from training data, bias amplification, job displacement concerns, misuse for scams or harmful content, and the difficulty of attributing accountability when AI-generated content causes harm.

SECTION 09

AI INTERVIEW SCENARIOS

Questions 186-195

Q186. Can you explain one of your AI or Machine Learning projects in detail?

Structure your answer using the STAR-style flow: (1) Problem/business context, (2) Dataset and preprocessing steps taken, (3) Models tried and why the final one was chosen, (4) Evaluation metrics and results achieved, and (5) Key challenges faced and how you solved them. Always tie the outcome back to measurable business or technical impact, e.g., 'improved churn prediction recall from 62% to 84% using SMOTE and XGBoost.'

Q187. How do you choose the right Machine Learning algorithm for a problem?

Consider the problem type (classification, regression, clustering), size and nature of the data, interpretability requirements, training/inference time constraints, and whether the relationship is likely linear or complex. It is best practice to start with a simple baseline model, then iterate to more complex models (ensembles, deep learning) only if needed.

Q188. What techniques would you use to improve model accuracy?

Techniques include better feature engineering, handling class imbalance (SMOTE, class weights), hyperparameter tuning (Grid/Random Search), trying ensemble methods (bagging/boosting), collecting more or cleaner data, and addressing overfitting/underfitting through regularization or model complexity adjustments.

Q189. How would you handle an imbalanced dataset?

An imbalanced dataset can be handled using resampling techniques like SMOTE (oversampling the minority class) or undersampling the majority class, adjusting class weights in the loss function, using appropriate metrics (F1, precision-recall AUC instead of accuracy), and trying algorithms robust to imbalance like tree-based ensembles.

Q190. How would you deploy a Machine Learning model into production?

A typical approach is to save/serialize the trained model (pickle, joblib, or ONNX), wrap it in a REST API using Flask/FastAPI, containerize it with Docker, deploy it on a scalable platform (Kubernetes, cloud services), set up CI/CD for updates, and add monitoring for drift and performance in production.

Q191. Can you explain an LLM or Generative AI project you have worked on?

Describe the use case (e.g., a RAG-based chatbot or FAQ assistant), how documents were chunked and embedded into a vector store, the retrieval and prompt design process, which LLM/API was used, and how you evaluated response quality/relevance, along with challenges like hallucination reduction or latency optimization.

Q192. How would you build an AI-powered chatbot from scratch?

Key steps: define the scope and intents, gather/prepare a knowledge base or FAQ dataset, choose an approach (rule-based, retrieval-based with RAG, or a fine-tuned/prompted LLM), build the backend pipeline (embeddings, vector store, LLM API integration), design the conversational flow, and deploy with a simple front-end interface, followed by testing and iteration based on user feedback.

Q193. How would you build a recommendation system for an e-commerce platform?

Approaches include collaborative filtering (using user-item interaction history), content-based filtering (using product attributes), or a hybrid model combining both. The pipeline typically involves collecting interaction data, building a user-item matrix or embeddings, training the recommendation model, and evaluating using metrics like precision@k or NDCG.

Q194. How would you debug and troubleshoot Machine Learning models?

Start by checking data quality and leakage, verify the train/validation/test splits are correct, examine learning curves for over/underfitting, inspect feature importance and residuals for unexpected patterns, and test with a simpler baseline model to isolate whether the issue is in data, features, or model choice.

Q195. How would you optimize the inference speed of a Deep Learning model?

Optimize inference speed through model quantization, pruning, knowledge distillation to a smaller model, converting to optimized formats like ONNX or TensorRT, batching requests, caching repeated queries, and deploying on appropriate hardware accelerators (GPU/TPU) suited to the workload.

SECTION 10

LATEST AI TRENDS

Questions 196-200

Q196. What are AI Agents, and how are they transforming AI applications?

AI Agents are LLM-powered systems that can autonomously plan, use tools, and take multi-step actions to accomplish a goal rather than just answering a single question. They are transforming AI applications by enabling automation of complex workflows -- like research, coding, and customer support -- that previously required significant human coordination.

Q197. What is Agentic AI, and how does it differ from traditional AI?

Agentic AI refers to systems designed with autonomy, planning, memory, and tool-use capability to pursue multi-step goals with minimal supervision. It differs from traditional AI, which typically performs a single, narrow, reactive task (like classifying an image) without ongoing planning or independent decision-making.

Q198. What is Multimodal AI, and what are its real-world applications?

Multimodal AI processes and reasons across multiple data types (text, image, audio, video) simultaneously. Real-world applications include visual question answering, AI assistants that can see and describe images, video summarization, accessibility tools for the visually impaired, and multimodal search engines.

Q199. What are Small Language Models (SLMs), and how do they compare with LLMs?

SLMs are compact language models (typically millions to a few billion parameters) optimized for efficiency. Compared to LLMs, they offer much faster inference, lower cost, and the ability to run on-device or at the edge, at the tradeoff of somewhat reduced general reasoning capability on very complex tasks.

Q200. What are the latest trends and future directions in Generative AI?

Current trends include increasingly capable AI agents that use tools autonomously, multimodal models that unify text/image/audio/video understanding, retrieval-augmented systems for grounded and up-to-date responses, smaller and more efficient on-device models, and growing focus on AI safety, alignment, and regulation as adoption accelerates across industries.

All the Best for Your AI / Machine Learning Interview!

This guide covers 200 questions across Artificial Intelligence, Machine Learning, Deep Learning, NLP & Generative AI, Python, Data Science, MLOps, and AI Ethics & Security.

Compiled and prepared by **THANGARAJ**
AI/ML Intern | Data Science Developer
Coimbatore, Tamil Nadu, India

Feel free to share this resource with anyone preparing for an AI or Data Science interview.